

CS343 Operating Systems

Lecture 13

Introduction to Memory Management



John Jose

Associate Professor

Department of Computer Science & Engineering



Indian Institute of Technology Guwahati

Background

- ❖ Program must be brought (from disk) into memory for a process to run
- ❖ Main memory and registers are only storage CPU can access directly
- ❖ Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- ❖ Register access in one CPU clock
- ❖ Main memory can take many cycles, causing a **stall**
- ❖ **Cache** sits between main memory and CPU registers
- ❖ Protection of memory required to ensure correct operation

Hardware Protection using Base and Limit Registers

- ❖ A pair of **base** and **limit registers** define the logical address space
- ❖ CPU must check every memory access generated in user mode to be sure it is between base and limit for that user



Address Binding

- ❖ Programs on disk, ready to be brought into memory to execute.
- ❖ Without support, must be loaded into address 0000
- ❖ Inconvenient to have user process physical address always at 0000
- ❖ Addresses represented in different ways in a program's life
 - ❖ Source code addresses usually symbolic
 - ❖ Compiled code addresses **bind** to relocatable addresses
 - ❖ i.e. "14 bytes from beginning of this module"
 - ❖ Linker or loader will bind relocatable addresses to absolute addresses
 - ❖ i.e. 7014

Binding of Instructions and Data to Memory

- ❖ Address binding of instructions and data to memory addresses can happen at three different stages
 - ❖ **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - ❖ **Load time:** Must generate **relocatable code** if memory location is not known at compile time
 - ❖ **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - ❖ Need hardware support for address maps (e.g., base and limit registers)

Logical vs. Physical Address Space


- ❖ **Logical address** – generated by the CPU; also referred to as **virtual address**
- ❖ **Physical address** – address seen by the memory unit
- ❖ Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- ❖ **Logical address space** is the set of all logical addresses generated by a program
- ❖ **Physical address space** is the set of all physical addresses allowed for a program

Memory-Management Unit (MMU)

- ❖ Hardware device that maps virtual to physical address at runtime
 - ❖ Base register now called **relocation register**
- ❖ The user program deals with logical addresses; it never sees the real physical addresses
 - ❖ Execution-time binding occurs when reference is made to location in memory
 - ❖ Logical address bound to physical addresses

Dynamic relocation using a relocation register


- ❖ Routine is not loaded until it is called
- ❖ Better memory-space utilization; unused routine is never loaded
- ❖ All routines kept on disk in relocatable load format
- ❖ Useful when large amounts of code are needed to handle infrequently occurring cases



Swapping

- ❖ A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
- ❖ Total virtual memory space of processes can exceed physical memory
- ❖ **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- ❖ **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- ❖ Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- ❖ System maintains a **ready queue** of ready-to-run processes which have memory images on disk

Swapping



Context Switch Time including Swapping


- ❖ If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- ❖ Context switch time can then be very high
- ❖ 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - ❖ Swap out time of 2 second
 - ❖ Plus swap in of same sized process
 - ❖ Total context switch swapping component time of 4000ms (4 seconds)
- ❖ Can reduce if reduce size of memory swapped – by knowing how much memory really being used
 - ❖ System calls to OS: `request_memory()` and `release_memory()`

Contiguous Allocation

- ❖ Main memory must support both OS and user processes
- ❖ Limited resource, must allocate efficiently
- ❖ Contiguous allocation is one early method
- ❖ Main memory has usually into two **partitions**:
 - ❖ Resident operating system, usually held in low memory with interrupt vector
 - ❖ User processes then held in high memory
 - ❖ Each process contained in single contiguous section of memory


Contiguous Allocation

- ❖ Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - ❖ Base register contains value of smallest physical address
 - ❖ Limit register contains range of logical addresses – each logical address must be less than the limit register
 - ❖ MMU maps logical address *dynamically*



Multiple-partition allocation

- ❖ Multiple-partition allocation
 - ❖ Degree of multiprogramming limited by number of partitions
 - ❖ **Variable-partition** sizes for efficiency as per size of process
 - ❖ **Hole** – block of available memory; holes of various size are scattered throughout memory
 - ❖ When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - ❖ Process exiting frees its partition, adjacent free partitions combined
 - ❖ Operating system maintains information about:
 - a) allocated partitions b) free partitions (hole)



Dynamic Storage-Allocation Problem

- ❖ How to satisfy a request of size n from a list of free holes?
- ❖ **First-fit**: Allocate the **first** hole that is big enough
- ❖ **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless ordered by size
 - ❖ Produces the smallest leftover hole
- ❖ **Worst-fit**: Allocate the **largest** hole; must also search entire list
 - ❖ Produces the largest leftover hole
- ❖ **Quick-fit**: Allocate the **first** hole that is big enough from the last point of allocation
- ❖ First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Fragmentation

- ❖ **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- ❖ **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- ❖ First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - ❖ $1/3$ may be unusable -> **50-percent rule**

Fragmentation

- ❖ Reduce external fragmentation by **compaction**
 - ❖ Shuffle memory contents to place all free memory together in one large block
 - ❖ Compaction is possible *only* if relocation is dynamic, and is done at execution time
 - ❖ I/O problem: Latch job in memory while it is involved in I/O



johnjose@iitg.ac.in
<http://www.iitg.ac.in/johnjose/>