

CS343 Operating Systems

Lecture 16

Virtual Memory Techniques



John Jose

Associate Professor

Department of Computer Science & Engineering

Indian Institute of Technology Guwahati



Structure of the Page Table

- ❖ Memory structures for paging can get huge using straight-forward methods
 - ❖ Hierarchical Paging
 - ❖ Hashed Page Tables
 - ❖ Inverted Page Tables

Hierarchical Page Tables

- ❖ Break up the logical address space into multiple page tables
- ❖ A simple technique is a two-level page table
- ❖ We then page the page table


Address-Translation Scheme



Hashed Page Tables

- ❖ The virtual page number is hashed into a page table
- ❖ This page table contains a chain of elements hashing to the same location
- ❖ Each element contains
 - (1) the virtual page number
 - (2) the value of the mapped page frame
 - (3) a pointer to the next element
- ❖ Virtual page numbers are compared in this chain searching for a match
- ❖ If a match is found, the corresponding physical frame is extracted


Hashed Page Tables



Inverted Page Table

- ❖ No page table per process
- ❖ One entry for each real page of memory
- ❖ Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- ❖ Decreases memory needed to store each page table,
- ❖ Increases time needed to search the table when a page reference occurs

Inverted Page Table Architecture



Shared Pages


❖ Shared code

- ❖ One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- ❖ Similar to multiple threads sharing the same process space
- ❖ Also useful for interprocess communication if sharing of read-write pages is allowed

❖ Private code and data

- ❖ Each process keeps a separate copy of the code and data
- ❖ The pages for the private code and data can appear anywhere in the logical address space

Shared Pages



Background


- ❖ Code needs to be in memory to execute
- ❖ But entire program code not needed at same time
- ❖ Consider ability to execute partially-loaded program
 - ❖ Program no longer constrained by limits of physical memory
 - ❖ Each program takes less memory while running → more programs run at the same time
 - ❖ Increased CPU utilization and throughput with no increase in response time or turnaround time

Background


- ❖ **Virtual memory** – separation of user logical memory from physical memory
- ❖ Only part of the program needs to be in memory for execution
- ❖ Logical address space can therefore be much larger than physical address space
- ❖ Allows address spaces to be shared by several processes
- ❖ More programs running concurrently
- ❖ Less I/O needed to load or swap processes

Background

- ❖ **Virtual address space** – logical view of how process is stored in memory
 - ❖ Usually start at address 0, contiguous addresses until end of space
 - ❖ Meanwhile, physical memory organized in page frames
 - ❖ MMU must map logical to physical
- ❖ Virtual memory can be implemented via Demand paging




Virtual Memory To Physical Memory Mapping




Demand Paging

- ❖ Bring a page into memory only when it is needed
 - ❖ Less I/O needed, no unnecessary I/O
 - ❖ Less memory needed
 - ❖ Faster response
 - ❖ More users
- ❖ **Lazy swapper** – never swaps a page into memory unless page will be needed




Valid-Invalid Bit



- ❖ With each page table entry a valid–invalid bit is associated (**v** $\Rightarrow$ in-memory – **memory resident**, **i** $\Rightarrow$ not-in-memory)
- ❖ Initially valid–invalid bit is set to **i** on all entries



logical memory



page table




Page Table When Some Pages Are Not in Main Memory

Page Fault

- ❖ If there is a reference to a page, first reference to that page will trap to operating system: **page fault** (page not found in main memory)
- ❖ Find free frame
- ❖ Swap page into frame via scheduled disk operation
- ❖ Reset tables to indicate page now in memory: Set validation bit = **1**
- ❖ Restart the instruction that caused the page fault

Steps in Handling a Page Fault



Demand Paging Overhead

❖ Page Fault Rate $0 \leq p \leq 1$

- if $p = 0$ no page faults
- if $p = 1$, every reference is a fault


❖ Effective Access Time (EAT)

$$\text{EAT} = (1 - p) \times \text{memory access} + p (\text{page fault overhead} \\ + \text{swap page out} + \text{swap page in})$$

Basic Page Replacement

- ❖ Find the location of the desired page on disk
- ❖ Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk if dirty
- ❖ Bring the desired page into the (newly) free frame; update the page and frame tables
- ❖ Continue the process by restarting the instruction that caused the trap

Page Replacement




Page Replacement Algorithm

❖ Page-replacement algorithm

- ❖ Want lowest page-fault rate on both first access and re-access
- ❖ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
 - ❖ String is just page numbers, not full addresses
 - ❖ Repeated access to the same page does not cause a page fault
 - ❖ **FIFO, LIFO,**
 - ❖ **Optimal,**
 - ❖ **LRU, LRU approximations,**
 - ❖ **LFU, MFU**

Page Faults Vs Number of Frames



First-In-First-Out (FIFO) Algorithm

- ❖ Ref. string: **7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1**
- ❖ 3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																
	0	0	0		2	2	4	4	4	0			0	0			7	7	7
					3	3	3	2	2	2			1	1			1	0	0
		1	1		1	0	0	0	3	3			3	2			2	2	1

page frames


15 page faults

- ❖ How to track ages of pages? - Use a FIFO queue

Belady's Anomaly

❖ Consider 1,2,3,4,1,2,5,1,2,3,4,5

❖ Adding more frames can cause more page faults! - **Belady's Anomaly**



Optimal Algorithm

- ❖ Replace page that will not be used for longest period of time
 - ❖ 9 is optimal for the example
- ❖ Practical difficulty- Can't read the future
- ❖ Used for measuring how well your algorithm performs

reference string


7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2				2				2			2				7
	0	0	0				0				0			0				0
		1	1				3				3			1				1

page frames

Least Recently Used (LRU) Algorithm

- ❖ Use past knowledge rather than future
- ❖ Replace page that has not been used in the most amount of time
- ❖ Associate time of last use with each page



- ❖ 12 faults – better than FIFO but worse than OPT
- ❖ Generally good algorithm and frequently used

LRU Algorithm Implementation

❖ Counter implementation

- ❖ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- ❖ When a page needs to be changed, look at the counters to find smallest value.

LRU Approximation Algorithms


❖ Reference bit

- ❖ With each page associate a bit, initially = 0
- ❖ When page is referenced, bit set to 1
- ❖ Replace any with reference bit = 0 (if one exists)


❖ Second-chance algorithm

- ❖ If page to be replaced has
 - ❖ Reference bit = 0 → replace it
 - ❖ reference bit = 1 then:
 - ❖ set reference bit 0, leave page in memory
 - ❖ replace next page, subject to same rules

Second-Chance Page-Replacement



(a)



(b)

Enhanced Second-Chance Algorithm

- ❖ Improve algorithm by using reference bit and modify bit
- ❖ Take ordered pair (reference, modify)
- ❖ (0, 0) neither recently used nor modified – best page to replace
- ❖ (0, 1) not recently used but modified – not quite as good, must write out before replacement
- ❖ (1, 0) recently used but clean – probably will be used again soon
- ❖ (1, 1) recently used and modified – probably will be used again soon and need to write out before replacement

Counting Algorithms

- ❖ Keep a counter of the number of references that have been made to each
- ❖ **Least Frequently Used (LFU) Algorithm**: replaces page with smallest count
- ❖ **Most Frequently Used (MFU) Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Page-Buffering Algorithms

- ❖ Always keep a pool of free frames
 - ❖ When needed, frame is always available, not found at fault time
 - ❖ Read page into free frame and select victim to evict and add to free pool
 - ❖ When convenient, evict victim
- ❖ Keep list of modified pages
 - ❖ When free, write to backing store and set to non-dirty
- ❖ Possibly, keep free frame contents intact and note what is in them
 - ❖ If referenced again before reused, no need to load contents again from disk
 - ❖ Generally useful to reduce penalty if wrong victim frame selected



johnjose@iitg.ac.in
<http://www.iitg.ac.in/johnjose/>