

CS 331 Tutorial #8

Functional Programming in Haskell



Introduction to Haskell



The Haskell Enigma

A Different Kind of Language

Haskell is a **purely functional programming language** — one that fundamentally challenges the way most developers are trained to think about code.

Rather than instructing the computer *how* to do something step by step, Haskell asks you to describe *what things are*. It's a profound shift in perspective.

Pure: No side effects in functions, Immutability is enforced, Effects are controlled

Functional Programming

A **functional programming** language is one where:

- Functions are first-class citizens (you can pass them around like data)
- Computation is expressed as evaluation of functions
- Emphasis on immutability and higher-order functions
- Examples: Lisp, Scala, Python (partially), JavaScript (partially)
- However, these languages allow side effects and mutable state.

Imperative vs. Functional: A Fundamental Shift

These two paradigms represent genuinely different philosophies about what a programme *is* and how it should be written.

Imperative Programming

You tell the computer a precise **sequence of tasks** to execute one after another. Variables hold state that can change at any moment — a variable assigned the value 5 today might be 10 tomorrow.

This is the dominant style in languages like C, Python, and Java. It mirrors how a machine actually operates: step by step, mutation by mutation.

λ Functional Programming

You define what things **are** using mathematical functions. State is **immutable** — once a value is bound to a name, it stays that way forever. There are no side effects hiding in the shadows.

This mirrors how mathematics works. Functions always map the same inputs to the same outputs, making code dramatically easier to reason about.

The Power of "What Is": Defining Reality with Functions

Instead of writing loops and tracking variables through mutations, Haskell lets you express ideas the way a mathematician would — concisely, elegantly, and without ambiguity.

These definitions are not just readable; they are **executable specifications**. The code and the concept are one and the same.

📄 This declarative style is one of Haskell's greatest strengths — your code says *what* to compute, not *how* to compute it.

Classic Functional Definitions

Factorial

The factorial of a number is the product of all integers from 1 up to that number. `factorial 0 = 1; factorial n = n * factorial (n-1)`

Sum of a List

The sum of a list is the first element plus the sum of the remainder. `sum [] = 0; sum (x:xs) = x + sum xs`

No Liars Allowed: The Immutability of Variables

Variables Are Truths

In Haskell, if you declare `a = 5`, then `a` is 5. Not "currently 5" — permanently, irrevocably 5. There is no reassignment, no mutation, no take-backs.

Fewer Bugs by Design

A huge category of bugs in imperative code stems from unexpected mutations — a value changes in one place and breaks logic somewhere else. Immutability makes this class of error **impossible**.

Predictable Code

When values never change, you can reason about any part of your programme in isolation. No hidden state. No surprise. What you see is what you get — always.



Referential Transparency: The Compiler's Best Friend

The Core Guarantee


A function in Haskell called with the **same parameters** will **always** return the same result — no matter when it is called, how many times, or in what order.

This property is known as **referential transparency**, and it is one of the most powerful guarantees in all of programming.

Why It Matters

- **Compiler optimisations** — the compiler can safely reorder, cache, or eliminate function calls it knows to be equivalent.
- **Easier reasoning** — developers can substitute any function call with its result mentally, making logic far easier to trace.
- **Safe parallelism** — with no shared mutable state, functions can be run in parallel without risk of race conditions.

Referential Transparency: Predictable & Reliable



Referential transparency is not merely a theoretical nicety — it is the foundation upon which Haskell's correctness guarantees are built. When every function behaves like a mathematical mapping from inputs to outputs, your entire programme becomes a system you can **trust**. The compiler leverages this to produce highly optimised code, and you leverage it to write tests that never need mocks or stubs.

Referential Transparency:
No Side Effect

Python

```
x = 10

def add_five():
    return x + 5
x= 10
print(add_five())
X = 20
print(add_five())
```

Output: 15 25
(depends on **global variable x**)

Haskell

```
addFive :: Int -> Int
addFive x = x + 5

main :: IO ()
main = do
    print (addFive 10)
    print (addFive 10)
```

Output: 15 15

Immutability

```
x = 10
print(x)

x = x + 5
print(x)
```

Python (Mutable variables)

```
main :: IO ()
main = do
    let x = 10
    let y = x + 5
    print x
    print y
```

Output: 10 15

`x = x + 1` ❌

Controlled Effects (I/O)

```
num = int(input("Enter a
number: "))
print(num + 5)
```

Input/output mixed directly with
logic

```
main :: IO ()
main = do
    putStrLn "Enter a number:"
    input <- getLine
    let num = read input :: Int
    print (num + 5)
```

Getting set up!

Installing Haskell

<https://www.haskell.org/ghcup/>



Youtube link : Installing Haskell on Windows (via ghcup)

<https://www.youtube.com/watch?v=bB4fmQiUYPw>



Youtube link : Installing Haskell on Ubuntu/Linux (via ghcup)


<https://www.youtube.com/watch?v=CUAZELnYwFs>



IDEs for Haskell

<https://wiki.haskell.org/IDEs>

Haskell: Values and Types





Expressions and Values

In Haskell, all computation is centred on the evaluation of expressions. Every expression, when evaluated, reduces to a value. There are no statements in the traditional imperative sense — instead, the language is built entirely around expressions and the values they produce.

Expression

A syntactic construct that can be evaluated. May be simple or complex, atomic or composed.

Value

The result of evaluating an expression. Values are the final, irreducible outputs of computation in Haskell.

The Concept of Type

Every value in Haskell is associated with a **type**. A type can be understood as a *set of values* — it defines what kind of data a value is, and constrains what operations may be applied to it. This association is fundamental to how Haskell programs are structured and verified.

Types as Sets

A type groups together values that share the same structure and permitted operations. For example, the type **Integer** contains all whole numbers.

Every Value Has a Type

No value in Haskell exists without a type. This universal association underpins the language's safety guarantees and reasoning model.

Types Constrain Behaviour

Knowing the type of a value tells you what you can and cannot do with it — preventing entire classes of programming errors before the program ever runs.

Examples of Expressions and Values

Haskell supports a rich variety of expression forms. Below are representative examples ranging from atomic literals to structured data and functions.

Atomic Values

`5` — an integer literal

`'a'` — a character literal

These are the simplest, indivisible values in the language.

Functions

`\x -> x + 1`

A lambda expression defining an anonymous function that increments its argument by one.

Structured Values

`[1, 2, 3]` — a list of integers

`('b', 4)` — a pair (tuple) of a character and an integer

Type Expressions

Just as expressions denote values, **type expressions denote types**. Haskell provides a concise notation for writing types, which appears in type signatures and other declarations throughout the language.

Type Expression	Description	Example Value
<code>Integer</code>	Infinite-precision whole numbers	<code>42</code>
<code>Char</code>	Single Unicode characters	<code>'z'</code>
<code>Integer->Integer</code>	Functions from integers to integers	<code>\x -> x+1</code>
<code>[Integer]</code>	Lists of integers	<code>[1,2,3]</code>
<code>(Char, Integer)</code>	Pairs of a character and an integer	<code>('a', 7)</code>

First-Class Values

In Haskell, **all values are first-class**. This is one of the most significant properties of a functional language. A first-class value is one that can be used in any context where a value is expected — without restriction.



Passed as Arguments

Functions can receive other functions or any value as input, enabling higher-order programming patterns.




Returned as Results

A function may return another function as its output, which is the basis for currying and partial application.



Stored in Data Structures

Values of any type, including functions, can be placed inside lists, tuples, or other data structures.



Typing: Associating Values with Types

The formal association of a value with its type is called a **typing**. In Haskell, the double colon notation `::` is used to express this association. Typings may appear in code as annotations or declarations, and are a core part of the language's specification.

1

`5 :: Integer`

The value `5` belongs to the type `Integer`.
This is the standard typing for a whole-number literal.

2

`'a' :: Char`

The character literal `'a'` is typed as `Char`,
the type of single Unicode characters.

3

`inc :: Integer->Integer`

The function `inc` takes an `Integer` and
returns an `Integer` — its type reflects its
input and output.

Declarations: Equations and Type Signatures

Haskell programs are built from **declarations**. The two most important forms are *equations*, which define the behaviour of functions, and *type signature declarations*, which explicitly state the type of a named value or function.

Equation (Definition)

Defines a function by specifying its behaviour:

```
inc n = n + 1
```

This binds the name **inc** to a function that adds 1 to its argument **n**.

Type Signature Declaration

States the type of a name explicitly:

```
inc :: Integer -> Integer
```

This declares that **inc** is a function from **Integer** to **Integer**, and appears alongside the equation in source code.

Type Safety and Error Detection

Haskell's **static type system** guarantees that well-typed programs are free from an entire class of runtime errors. Type checking occurs at *compile time*, so type mismatches are caught before the programme ever runs.

→ Type-Safe by Design

The type system enforces that operations are only applied to values of the correct type. An expression like `'a' + 'b'` is rejected by the compiler because addition is not defined for characters.

→ Compile-Time Detection

Type errors are identified during compilation, not at runtime. This eliminates a large category of bugs and makes programmes more predictable and robust.

→ Limits of the Type System

Not all errors are type errors. For example, `1/0` is type-correct (both operands are numeric) but will produce a runtime error. The type system cannot catch logical or arithmetic faults of this kind.

Types

Polymorphic types

A function works with **many types**, using type variables..

```
identity :: a -> a
identity x = x
```

lower case

```
main :: IO ()
main = do
    print (identity 10)
    print (identity
"Hello")
```

User-Defined Types

We can define our own types in Haskell using a data declaration

```
data Color = Red | Green
           | Blue
```

```
main :: IO ()
main = print Red
```

```
data Person = Person String Int
```

tuple type

Recursive Types

A type defined in terms of itself.

```
data Tree a = Leaf a |
             Branch (Tree a) (Tree a)
```

a polymorphic binary tree type whose elements are either leaf nodes containing a value of type `a`, or internal nodes ("branches") containing (recursively) two sub-trees.

Haskell provides a way to define **type synonyms**; i.e. names for commonly used types.

```
type Name = String
type Age = Int
```

```
data Person = Person Name Age
```



Functions in Haskell

Understanding Haskell Functions

The Functional Paradigm

In Haskell, functions are not merely a feature — they are the **primary mechanism** of computation and abstraction. Unlike imperative languages that rely on statements and mutable state, Haskell expresses all computation through the application and composition of functions.

Why Functions Matter

Because Haskell is a **purely functional language**, understanding how functions are defined, applied, and composed is foundational. Every program is, at its heart, a collection of function definitions and their interactions.

- Functions in Haskell are **pure**: given the same input, they always return the same output, with no side effects.



Defining Functions: The **add** Example

In Haskell, functions are called by writing the function name, a space and then the parameters, separated by spaces

Type Signature

In Haskell, a function definition begins with its **type signature**, which declares the types of its inputs and output. For the **add** function:

```
add :: Integer -> Integer -> Integer
add x y = x + y
```

This declares that **add** takes two **Integer** arguments and returns an **Integer** result — the sum of both inputs.

Reading the Definition

Type Signature

add :: Integer -> Integer -> Integer
specifies the function's contract with the type system.

Equation

add x y = x + y is the defining equation: pattern-match two arguments and return their sum.

Curried Functions: A Key Concept

A **curried function** is a function that takes **one argument at a time** and returns another function until all arguments are provided.

The definition `add x y = x + y` is a canonical example of a **curried function**. Currying — named after logician Haskell Curry — means that a function takes its arguments *one at a time*, returning a new function after each application.

1

Type Equivalence

`Integer -> Integer -> Integer` is right-associative, equivalent to `Integer -> (Integer -> Integer)`.

2

First Application

Applying `add` to one argument, e.g. `add 3`, does **not** compute a result — it returns a new function of type `Integer -> Integer`.

3

Second Application

Supplying the second argument, e.g. `add 3 5`, completes the evaluation and returns the `Integer` value `8`.

❏ All multi-argument functions in Haskell are curried by default. This is not optional syntax — it is the fundamental semantics of the language.

Partial Application: Building New Functions

Instead of `f(x,y)` Haskell does `(f x) y`

```
add :: Int -> Int -> Int
```

```
add x y = x + y
```

`add 5` returns a function

```
= (\y -> 5 + y)
```

```
= (add 5) 3
```

```
= 8
```

Practical Significance

1 Code Reuse

Specialise general functions into specific ones without rewriting logic, reducing duplication.

2 Expressiveness

Concisely describe computations — `inc` communicates intent more clearly than a full definition.

3 Function Values

Partial application confirms that functions are **values** in Haskell: they can be bound to names, stored, and passed around.

What Is Partial Application?

Partial application is the technique of supplying fewer arguments to a curried function than it expects. Rather than producing a final value, this yields a **new function** with some arguments already fixed.

Curried function (default in Haskell)

```
add :: Int -> Int -> Int
add x y = x + y
```

```
addFive :: Int -> Int
addFive = add 5
```

```
main = print (addFive 10)
```

Output: 15

You fixed one argument → got a new function. This is **partial application**

```
main
= print (addFive 10)
= print ((add 5) 10)
= print ((\y -> 5 + y) 10)
= print (5 + 10)
= print 15
= 15
```

Tuple function (non-curried)

```
addTuple :: (Int, Int) -> Int
addTuple (x, y) = x + y
```

```
addFive = addTuple (5, ?)    --  invalid
```

Partial application (NOT possible directly)

You must provide both arguments at once

```
inc = add 1
```

Can define a new function 'inc' reusing add

Higher-Order Functions: Functions as Arguments and Return Values

A **higher-order function** is one that accepts functions as arguments, returns functions as results, or both. The canonical example in Haskell is **map**.

The **map** Function

```
map :: (a -> b) -> [a] -> [b]
map f []      = []
map f (x:xs) = f x : map f xs
```

map takes a function **f** and a list, and applies **f** to every element, collecting results into a new list. The type variable **a** and **b** make it **polymorphic** — it works on lists of any type.

Example

Input

```
map (add 1) [1, 2, 3]
```

Process

Apply **add 1** (a partially applied function) to each element: 1+1, 2+1, 3+1.

Output

```
[2, 3, 4] — a transformed list of the same length.
```

Lambda Abstractions: Anonymous Functions

Syntax and Semantics

Haskell supports **lambda abstractions** — anonymous functions written inline, without a named binding. The syntax mirrors the mathematical lambda calculus:

```
\x -> x + 1      -- equivalent to inc
\x -> \y -> x + y -- equivalent to add
\x y -> x + y    -- shorthand form
```

The general form `\x -> exp` has type `t1 -> t2` if variable `x` has type `t1` and expression `exp` has type `t2`.

Meaning:
`\x -> input`
`-> -> maps to`
`x + 1 -> result`

```
(\x -> x + 1)5
Output: 6
```

Application

When to Use Lambdas

→ Inline Arguments

Pass a small function directly to `map`, `filter`, or `foldr` without cluttering the namespace with a named definition.

→ Clarity of Intent


Make it explicit that a function is being constructed at that point, rather than referring to a definition elsewhere.

→ Equivalence to Named Definitions

Any lambda can be converted to a named function and vice versa — they are semantically identical.

Infix Operators: Functions in Disguise


In Haskell, **infix operators** such as `+`, `++`, and `.` are simply functions with symbolic names, defined using the same equation syntax as any other function.



List Concatenation (`++`)

```
(++)      :: [a] -> [a] -> [a]
[] ++ ys   = ys
(x:xs) ++ ys = x : (xs ++ ys)
```

Appends two lists by recursively prepending elements of the first list onto the second. The base case is the empty list.



Function Composition (`.`)

```
(.) :: (b->c) -> (a->b) -> (a->c)
f . g = \x -> f (g x)
```

Composes two functions: the output of `g` becomes the input of `f`. Essential for building clean, readable pipelines of transformations.

Sections: Partial Application of Infix Operators

What Are Sections?

Sections provide a concise syntax for partially applying an infix operator, producing a new function. They are enclosed in parentheses with one argument supplied and the other omitted.

```
(x +)  ≡  \y -> x + y  -- left section
(+ y)  ≡  \x -> x + y  -- right section
(+)     ≡  \x y -> x + y -- full operator
```

This mechanism allows infix operators to be used seamlessly wherever a function value is expected.

Practical Usage

1

Passing to **map**

map (+1) [1,2,3] increments every list element — a right section used as a higher-order argument.

2

Passing to **filter**

filter (>0) xs retains only positive elements — a right section forming a predicate function.

3

Combining with Composition

(+1) . (*2) doubles then increments — sections compose naturally with **(.)**.

Summary: Functions as First-Class Citizens



Currying

All functions are curried by default, enabling elegant partial application and specialisation.



Higher-Order Functions

Functions are values — they can be passed, returned, and composed, enabling powerful abstractions like `map`.



Lambda Abstractions


Anonymous functions allow inline definition, keeping code concise and rooted in mathematical notation.



Operators & Sections

Infix operators are functions, and sections make partial application of operators natural and readable.

Case Expressions and Pattern Matching in Haskell



Introduction to Pattern Matching

What is Pattern Matching?

Pattern matching is a core feature in Haskell for defining functions. It allows functions to be written as a series of equations, each covering a different structural form of the input data.

Rather than using conditionals to inspect values, Haskell lets you describe the *shape* of data directly in the function definition.

Key Distinction from Prolog

Haskell's pattern matching is strictly **one-way**: patterns are matched against values, not unified with them. There is no implicit backtracking as found in Prolog. Each equation is tried in order, top to bottom, and the first successful match is used.

This deterministic behaviour makes reasoning about Haskell programs straightforward.

Data Constructor Patterns

Patterns may be constructed using the constructors of any type — whether built-in or user-defined. This includes lists, tuples, strings, numbers, and characters. Crucially, patterns may be **nested to arbitrary depth**, allowing highly precise structural decomposition.

Example Function

```
contrived :: ([a], Char, (Int, Float), String, Bool) ->
```

```
Bool
```

```
contrived ([], 'b', (1, 2.0), "hi", True) = False
```

This single equation matches a 5-tuple whose components each conform to a specific pattern — an empty list, a literal character, a nested integer-float pair, a string, and a boolean.

Supported Pattern Forms

- Literal values: numbers, characters, strings
- Data constructors: **True**, **False**, **Nothing**, **Just x**
- List constructors: **[]**, **(x:xs)**
- Tuples: **(a, b)**, **(a, b, c)**
- User-defined constructors from **data** declarations

Example: if we wanted a function that says the numbers from 1 to 5 and says "Not between 1 and 5" for any other number? (with pattern matching and without if-then-else)

```
sayMe :: (Integral a) => a -> String

sayMe 1 = "One!"
sayMe 2 = "Two!"
sayMe 3 = "Three!"
sayMe 4 = "Four!"
sayMe 5 = "Five!"
sayMe x = "Not between 1 and 5"
```

This program in Haskell defines a custom data type and a function to compute the area of shapes

```
data Shape = Circle Float | Rectangle Float Float

area :: Shape -> Float
area (Circle r) = 3.14 * r * r
area (Rectangle w h) = w * h
```

Irrefutable vs. Refutable Patterns

Irrefutable Patterns

Always succeed in matching any value. The most common irrefutable pattern is a **formal parameter** (a plain variable name). When an irrefutable pattern is matched, it binds the variable to the value.

Patterns must be **linear**: no variable may appear more than once within a single equation. For example, **`f x x = ...`** is not permitted.

Refutable Patterns

May fail to match a value, causing Haskell to move on to the next equation. Data constructor patterns are the primary example — **`(x:xs)`** will fail if the argument is an empty list **`[]`**.

When all equations fail to match, a **Non-exhaustive patterns** runtime error is raised, making completeness an important design concern.

As-patterns

Purpose and Syntax

As-patterns allow you to give a name to an entire pattern so it can be referenced on the right-hand side of an equation, without repeating the pattern itself. You do that by putting a name and an @ in front of a pattern.

For example, a function that duplicates the first element in a list might be written as:

```
f (x:xs) = x:x:xs
```

```
f s@(x:xs) = x : s
```

Here, **s** is bound to the whole list **(x:xs)**. Both the name **s** and the components **x** and **xs** are available on the right-hand side.

Why Use As-patterns?

- Avoids reconstructing a data structure you already have
- Improves readability by giving the whole structure a meaningful name
- Can reduce memory allocation by reusing the original structure
- As-patterns always result in a successful match for the named variable, even if the sub-pattern subsequently fails

```
capital :: String -> String
capital "" = "Empty string, whoops!"
capital all@(x:xs) = "The first letter of " ++ all ++ "
is " ++ [x]
```

```
ghci> capital "Institute"
"The first letter of Institute is I"
```

Wild-cards ()

The wild-card pattern is used to match values that are not needed by the function's logic. It explicitly signals that a portion of the input is intentionally ignored, which improves code clarity.



Matches Anything, Binds Nothing

Unlike a variable, a wild-card does not introduce a binding. You cannot reference on the right-hand side. Each occurrence is independent and matches any value silently.



Multiple Wild-cards Allowed

A single equation may contain multiple wild-cards.



Communicates Intent

Using rather than an unused variable name is idiomatic Haskell. It advertises that the omission is deliberate, and GHC will warn you if a named variable goes unused — the wild-card suppresses this warning.

. For example: `head (x:_)` = `x` and `tail (_:xs)` = `xs`

Pattern Matching Semantics

Understanding the semantics of pattern matching means understanding how Haskell decides *which* equation to use, and what happens at each step of the matching process.

1

Order of Attempts

Equations are tried strictly from top to bottom. The first equation whose pattern matches the argument is selected.

2

Handling Failures

If a refutable pattern fails, Haskell proceeds to the next equation. If no equation matches, a runtime error is raised.

3

Refutability

Each individual pattern is either irrefutable (always succeeds) or refutable (may fail), shaping the overall flow of evaluation.

Guards

Whereas patterns are a way of making sure a value conforms to some form and deconstructing it, guards are a way of testing whether some property of a value (or several of them) are true or false.

Guards are used when decisions depend on conditions (like comparisons), not just structure.

```
bmiTell :: Float -> String
bmiTell bmi
  | bmi <= 18.5 = "Underweight"
  | bmi <= 25.0 = "Normal"
  | otherwise  = "Overweight"
```

Guards are indicated by pipes that follow a function's name and its parameters. The first is basically a boolean expression. If it evaluates to **True**, then the corresponding function body is used. If it evaluates to **False**, checking drops through to the next guard and so on.

Many times, the last guard is **otherwise**. **otherwise** is defined simply as **otherwise = True** and catches everything.

Where bindings

where is used to define local variables or helper functions that are used in a function.

It helps make code, cleaner, more readable, and less repetitive.

Variables defined in **where**, are visible only inside that function, not accessible outside

```
bmiTell :: Float -> Float -> String
bmiTell weight height
  | weight / (height * height) <= 18.5 = "Underweight"
  | weight / (height * height) <= 25.0 = "Normal"
  | otherwise = "Overweight"
```

```
area :: Float -> Float
area r = piVal * r * r
  where
    piVal = 3.14
```

```
bmiTell :: Float -> Float -> String
bmiTell weight height
  | bmi <= 18.5 = "Underweight"
  | bmi <= 25.0 = "Normal"
  | otherwise  = "Overweight"
  where
    bmi = weight / (height * height)
```

Let bindings

let is used to define local variables or functions inside an expression

where bindings are a syntactic construct that let you bind to variables at the end of a function and the whole function can see them, including all the guards.

let you bind to variables anywhere and are expressions themselves, but are very local, they don't span across guards

```
main = print (  
    let x = 3  
        y = 4  
    in x + y  
)
```

```
calc :: Int -> Int  
calc x =  
    let expr = x * x + 1  
    in expr + expr
```

let → define something locally and use it immediately
where → define something for the whole function

Case Expressions

A **case expression** in Haskell is used to perform pattern matching on a value within an expression and select the appropriate result based on matching patterns.

Case version

```
describe :: [a] -> String
describe xs =
  case xs of
    [] -> "Empty"
    (x:_) -> "Non-empty"
```

Pattern matching version

```
describe :: [a] -> String
describe [] = "Empty"
describe (x:_) = "Non-empty"
```

Formal Parameters as Patterns

Variables are Patterns

In Haskell, formal parameters in function definitions are themselves a special case of pattern: the **variable pattern**. A variable pattern never fails to match — it accepts any value and binds that value to the variable name for use in the function body.

This universality is what makes variable patterns irrefutable. The consequence is the **linearity requirement**: within a single equation, each variable name may appear at most once on the left-hand side. Repeating a variable to enforce equality is not supported in standard pattern syntax.

Linearity Rule

The following is **not** valid Haskell:

```
equal x x = True
```

Instead, use a guard:

```
equal x y | x == y = True | otherwise = False
```

Or use the built-in `(==)` operator directly.

Conclusion: The Power of Pattern Matching

Pattern matching is not merely syntactic convenience — it is a foundational design principle of idiomatic Haskell. It enables function definitions that are concise, readable, and structurally explicit.

Conciseness

Eliminates verbose conditional chains. Each equation handles a distinct structural case cleanly and directly.

Expressiveness

Allows deep deconstruction of arbitrarily nested data structures in a single, readable left-hand side.

Idiomatic Style

Pattern matching is central to Haskell's identity. Mastering it — including as-patterns, wild-cards, and refutability — is essential to writing effective Haskell code.

